

In [1]:

```

!pip install --upgrade keras_applications==1.0.6 keras==2.2.4

import os
if os.path.exists('./data') == False:
    from modelarts.session import Session
    session = Session()

    session.download_data(
        bucket_path="modelarts-labs/notebook/DL_image_recognition/image_recognition.tar.gz",
        path="./image_recognition.tar.gz")

    # 使用tar命令解压资源包
    !tar xf ./image_recognition.tar.gz

    # 清理压缩包
    !rm -f ./image_recognition.tar.gz

```

Requirement already up-to-date: keras_applications==1.0.6 in /opt/conda/envs/python36_tf/lib/python3.6/site-packages

Requirement already up-to-date: keras==2.2.4 in /opt/conda/envs/python36_tf/lib/python3.6/site-packages

Requirement already up-to-date: h5py in /opt/conda/envs/python36_tf/lib/python3.6/site-packages (from keras_applications==1.0.6)

Requirement already up-to-date: numpy>=1.9.1 in /opt/conda/envs/python36_tf/lib/python3.6/site-packages (from keras_applications==1.0.6)

Requirement already up-to-date: pyyaml in /opt/conda/envs/python36_tf/lib/python3.6/site-packages (from keras==2.2.4)

Requirement already up-to-date: keras-preprocessing>=1.0.5 in /opt/conda/envs/python36_tf/lib/python3.6/site-packages (from keras==2.2.4)

Requirement already up-to-date: six>=1.9.0 in /home/jovyan/modelarts-sdk (from keras==2.2.4)

Requirement already up-to-date: scipy>=0.14 in /opt/conda/envs/python36_tf/lib/python3.6/site-packages (from keras==2.2.4)

You are using pip version 9.0.1, however version 19.1.1 is available.

You should consider upgrading via the 'pip install --upgrade pip' command.

In [2]:

```

from keras.applications.vgg16 import VGG16
from keras.preprocessing import image
import numpy as np

from keras.preprocessing import image
from keras.models import Model
from keras.layers import Dense, GlobalAveragePooling2D
from keras import backend as K
from keras.models import load_model

from keras.preprocessing.image import ImageDataGenerator

```

Using TensorFlow backend.

In [3]:

```
import os
from PIL import Image
def load_data():
    dirname = "./data"
    path = "./data"

    num_train_samples = 25000

    x_train = np.empty((num_train_samples, 224, 224, 3), dtype='uint8')
    y_train = np.empty((num_train_samples, 1), dtype='uint8')
    index = 0
    for file in os.listdir("./data"):
        image = Image.open(os.path.join(dirname, file)).resize((224, 224))
        image = np.array(image)
        x_train[index, :, :, :] = image

        if "cat" in file:
            y_train[index, 0] = 1
        elif "dog" in file:
            y_train[index, 0] = 0

        index += 1
    return (x_train, y_train)
```

In [4]:

```
(x_train, y_train) = load_data()
```

In [5]:

```
print(x_train.shape)
print(y_train.shape)
```

```
(25000, 224, 224, 3)
(25000, 1)
```

In [6]:

```
from keras.utils import np_utils
def process_data(x_train, y_train):
    x_train = x_train.astype(np.float32)
    x_train /= 255
    n_classes = 2
    y_train = np_utils.to_categorical(y_train, n_classes)
    return x_train, y_train
```

In [7]:

```
def build_model(base_model):
    x = base_model.output
    x = GlobalAveragePooling2D()(x)
    predictions = Dense(2, activation='softmax')(x)
    model = Model(inputs=base_model.input, outputs=predictions)
    print(type(model))
    return model
```

In [8]:

```
x_train, y_train= process_data(x_train, y_train)
print(x_train.shape)
print(y_train.shape)
```

```
(25000, 224, 224, 3)
```

```
(25000, 2)
```

In [9]:

```
base_model = VGG16(weights='imagenet', include_top=False)
```

Downloading data from https://github.com/fchollet/deep-learning-models/releases/download/v0.1/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5 (https://github.com/fchollet/deep-learning-models/releases/download/v0.1/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5)

```
58892288/58889256 [=====] - 222s 4us/step
```

In [10]:

```
for layer in base_model.layers:
    layer.trainable = False

model = build_model(base_model)
model.summary()
```

<class 'keras.engine.training.Model'>

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, None, None, 3)	0
block1_conv1 (Conv2D)	(None, None, None, 64)	1792
block1_conv2 (Conv2D)	(None, None, None, 64)	36928
block1_pool (MaxPooling2D)	(None, None, None, 64)	0
block2_conv1 (Conv2D)	(None, None, None, 128)	73856
block2_conv2 (Conv2D)	(None, None, None, 128)	147584
block2_pool (MaxPooling2D)	(None, None, None, 128)	0
block3_conv1 (Conv2D)	(None, None, None, 256)	295168
block3_conv2 (Conv2D)	(None, None, None, 256)	590080
block3_conv3 (Conv2D)	(None, None, None, 256)	590080
block3_pool (MaxPooling2D)	(None, None, None, 256)	0
block4_conv1 (Conv2D)	(None, None, None, 512)	1180160
block4_conv2 (Conv2D)	(None, None, None, 512)	2359808
block4_conv3 (Conv2D)	(None, None, None, 512)	2359808
block4_pool (MaxPooling2D)	(None, None, None, 512)	0
block5_conv1 (Conv2D)	(None, None, None, 512)	2359808
block5_conv2 (Conv2D)	(None, None, None, 512)	2359808
block5_conv3 (Conv2D)	(None, None, None, 512)	2359808
block5_pool (MaxPooling2D)	(None, None, None, 512)	0
global_average_pooling2d_1 ((None, 512)	0
dense_1 (Dense)	(None, 2)	1026
Total params: 14,715,714		
Trainable params: 1,026		
Non-trainable params: 14,714,688		

In [11]:

```
import keras
opt = keras.optimizers.rmsprop(lr=0.0001, decay=1e-6)
model.compile(loss='binary_crossentropy',
              optimizer=opt,
              metrics=['accuracy'])
```

In [12]:

```
from keras.callbacks import ModelCheckpoint, EarlyStopping, ReduceLRonPlateau
es = EarlyStopping(monitor='val_acc', baseline=0.9, patience=15, verbose=1, mode='auto')
cp = ModelCheckpoint(filepath="./ckpt_vgg16_dog_and_cat.h5", monitor="val_acc", verbose=1, save_best_only=True)
lr = ReduceLRonPlateau(monitor="val_acc", factor=0.1, patience=10, verbose=1, mode="auto", min_lr=0)
callbacks = [es, cp, lr]
```

In [13]:

```
history = model.fit(x=x_train,
                   y=y_train,
                   batch_size=16,
                   epochs=100,
                   verbose=1,
                   callbacks=callbacks,
                   validation_split=0.25,
                   shuffle=True,
                   initial_epoch=0)
```

```
18750/18750 [=====  
0.9277 - val_loss: 0.1932 - val_acc: 0.9214
```

Epoch 00096: val_acc did not improve from 0.92240

Epoch 97/100

```
18750/18750 [=====  
0.9277 - val_loss: 0.1932 - val_acc: 0.9214
```

Epoch 00097: val_acc did not improve from 0.92240

Epoch 98/100

```
18750/18750 [=====  
0.9277 - val_loss: 0.1932 - val_acc: 0.9214
```

Epoch 00098: val_acc did not improve from 0.92240

Epoch 99/100

```
18750/18750 [=====  
0.9277 - val_loss: 0.1932 - val_acc: 0.9216
```


Epoch 00099: val_acc did not improve from 0.92240

Epoch 00099: early stopping

In [15]:


```
import matplotlib.pyplot as plt
print(history.history['val_acc'][0])
# 绘制训练 & 验证的准确率值
plt.plot(history.history['acc'])
plt.plot(history.history['val_acc'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()
```

0.8382400000190735



In [16]:

```
# 绘制训练 & 验证的损失值
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()
```



In [17]:

```
base_model = VGG16(weights='imagenet', include_top=False)
model = build_model(base_model)
model.compile(loss='binary_crossentropy',
              optimizer=opt,
              metrics=['accuracy'])
model.summary()
```

<class 'keras.engine.training.Model'>

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	(None, None, None, 3)	0
block1_conv1 (Conv2D)	(None, None, None, 64)	1792
block1_conv2 (Conv2D)	(None, None, None, 64)	36928
block1_pool (MaxPooling2D)	(None, None, None, 64)	0
block2_conv1 (Conv2D)	(None, None, None, 128)	73856
block2_conv2 (Conv2D)	(None, None, None, 128)	147584
block2_pool (MaxPooling2D)	(None, None, None, 128)	0
block3_conv1 (Conv2D)	(None, None, None, 256)	295168
block3_conv2 (Conv2D)	(None, None, None, 256)	590080
block3_conv3 (Conv2D)	(None, None, None, 256)	590080
block3_pool (MaxPooling2D)	(None, None, None, 256)	0
block4_conv1 (Conv2D)	(None, None, None, 512)	1180160
block4_conv2 (Conv2D)	(None, None, None, 512)	2359808
block4_conv3 (Conv2D)	(None, None, None, 512)	2359808
block4_pool (MaxPooling2D)	(None, None, None, 512)	0
block5_conv1 (Conv2D)	(None, None, None, 512)	2359808
block5_conv2 (Conv2D)	(None, None, None, 512)	2359808
block5_conv3 (Conv2D)	(None, None, None, 512)	2359808
block5_pool (MaxPooling2D)	(None, None, None, 512)	0
global_average_pooling2d_2 ((None, 512)		0
dense_2 (Dense)	(None, 2)	1026
Total params: 14,715,714		
Trainable params: 14,715,714		
Non-trainable params: 0		

In [18]:

```
history = model.fit(x=x_train,
                    y=y_train,
                    batch_size=16,
                    epochs=5,
                    verbose=1,
                    callbacks=callbacks,
                    validation_split=0.25,
                    shuffle=True,
                    initial_epoch=0)
```

Train on 18750 samples, validate on 6250 samples

Epoch 1/5

```
18750/18750 [=====] - 168s 9ms/step - loss: 0.3699 - acc:
0.8477 - val_loss: 0.1915 - val_acc: 0.9246
```

Epoch 00001: val_acc improved from 0.92240 to 0.92464, saving model to ./ckpt_vgg16_dog_and_cat.h5

Epoch 2/5

```
18750/18750 [=====] - 166s 9ms/step - loss: 0.1488 - acc:
0.9441 - val_loss: 0.1298 - val_acc: 0.9496
```

Epoch 00002: val_acc improved from 0.92464 to 0.94960, saving model to ./ckpt_vgg16_dog_and_cat.h5

Epoch 3/5

```
18750/18750 [=====] - 164s 9ms/step - loss: 0.1103 - acc:
0.9570 - val_loss: 0.1028 - val_acc: 0.9598
```

Epoch 00003: val_acc improved from 0.94960 to 0.95984, saving model to ./ckpt_vgg16_dog_and_cat.h5

Epoch 4/5

```
18750/18750 [=====] - 164s 9ms/step - loss: 0.0910 - acc:
0.9651 - val_loss: 0.0892 - val_acc: 0.9654
```

Epoch 00004: val_acc improved from 0.95984 to 0.96544, saving model to ./ckpt_vgg16_dog_and_cat.h5

Epoch 5/5


```
18750/18750 [=====] - 164s 9ms/step - loss: 0.0778 - acc:
0.9713 - val_loss: 0.0850 - val_acc: 0.9677
```

Epoch 00005: val_acc improved from 0.96544 to 0.96768, saving model to ./ckpt_vgg16_dog_and_cat.h5

In [19]:

```
import matplotlib.pyplot as plt
print(history.history['val_acc'][0])
# 绘制训练 & 验证的准确率值
plt.plot(history.history['acc'])
plt.plot(history.history['val_acc'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show();
```

0.92464



In []: