

In [2]:

```
import os
if os.path.exists('./data') == False:
    from modelarts.session import Session
    session = Session()

    session.download_data(
        bucket_path="modelarts-labs/end2end/image_recognition/dog_and_cat_25000.tar.gz",
        path="./dog_and_cat_25000.tar.gz")

    # 使用tar命令解压资源包
    !tar xf ./dog_and_cat_25000.tar.gz

    # 清理压缩包
    !rm -f ./dog_and_cat_25000.tar.gz
```

Successfully download file modelarts-labs/end2end/image_recognition/dog_and_cat_25000.tar.gz from OBS to local ./dog_and_cat_25000.tar.gz

In [3]:

```
from keras.applications.vgg16 import VGG16
from keras.preprocessing import image
from keras.applications.resnet50 import preprocess_input, decode_predictions
import numpy as np

from keras.applications.mobilenetv2 import MobileNetV2
from keras.preprocessing import image
from keras.models import Model
from keras.layers import Dense, GlobalAveragePooling2D
from keras import backend as K
from keras.models import load_model

from keras.preprocessing.image import ImageDataGenerator
import os
from PIL import Image
```

Using TensorFlow backend.

In [4]:

```
def load_data():
    dirname = "./data"
    path = "./data"

    num_train_samples = 25000

    x_train = np.empty((num_train_samples, 224, 224, 3), dtype='uint8')
    y_train = np.empty((num_train_samples, 1), dtype='uint8')
    index = 0
    for file in os.listdir("./data"):
        image = Image.open(os.path.join(dirname, file)).resize((224, 224))
        image = np.array(image)
        x_train[index, :, :, :] = image

        if "cat" in file:
            y_train[index, 0] = 1
        elif "dog" in file:
            y_train[index, 0] = 0

        index += 1
    return (x_train, y_train)
```

In [5]:

```
(x_train, y_train) = load_data()
print(x_train.shape)
print(y_train.shape)
```

```
(25000, 224, 224, 3)
(25000, 1)
```

In [6]:

```
from keras.utils import np_utils
def process_data(x_train, y_train):
    x_train = x_train.astype(np.float32)
    x_train /= 255
    n_classes = 2
    y_train = np_utils.to_categorical(y_train, n_classes)
    return x_train, y_train
```

In [7]:

```
x_train, y_train = process_data(x_train, y_train)
print(x_train.shape)
print(y_train.shape)
```

```
(25000, 224, 224, 3)
(25000, 2)
```

In [8]:

```
def build_model(base_model):  
    x = base_model.output  
    x = GlobalAveragePooling2D()(x)  
    predictions = Dense(2, activation='softmax')(x)  
    model = Model(inputs=base_model.input, outputs=predictions)  
    print(type(model))  
    return model
```

In [9]:

```
base_model = VGG16(weights=None, include_top=False)
```

In [10]:

```
model = build_model(base_model)
model.summary()
```

```
<class 'keras.engine.training.Model'>
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, None, None, 3)	0
block1_conv1 (Conv2D)	(None, None, None, 64)	1792
block1_conv2 (Conv2D)	(None, None, None, 64)	36928
block1_pool (MaxPooling2D)	(None, None, None, 64)	0
block2_conv1 (Conv2D)	(None, None, None, 128)	73856
block2_conv2 (Conv2D)	(None, None, None, 128)	147584
block2_pool (MaxPooling2D)	(None, None, None, 128)	0
block3_conv1 (Conv2D)	(None, None, None, 256)	295168
block3_conv2 (Conv2D)	(None, None, None, 256)	590080
block3_conv3 (Conv2D)	(None, None, None, 256)	590080
block3_pool (MaxPooling2D)	(None, None, None, 256)	0
block4_conv1 (Conv2D)	(None, None, None, 512)	1180160
block4_conv2 (Conv2D)	(None, None, None, 512)	2359808
block4_conv3 (Conv2D)	(None, None, None, 512)	2359808
block4_pool (MaxPooling2D)	(None, None, None, 512)	0
block5_conv1 (Conv2D)	(None, None, None, 512)	2359808
block5_conv2 (Conv2D)	(None, None, None, 512)	2359808
block5_conv3 (Conv2D)	(None, None, None, 512)	2359808
block5_pool (MaxPooling2D)	(None, None, None, 512)	0
global_average_pooling2d_1 ((None, 512)		0
dense_1 (Dense)	(None, 2)	1026
Total params: 14,715,714		
Trainable params: 14,715,714		
Non-trainable params: 0		

In [11]:

```
import keras
opt = keras.optimizers.rmsprop(lr=0.0001, decay=1e-6)
model.compile(loss='categorical_crossentropy',
              optimizer=opt,
              metrics=['accuracy'])
```

In [12]:

```
from keras.callbacks import ModelCheckpoint, EarlyStopping
es = EarlyStopping(monitor='val_acc', baseline=0.9, patience=30, verbose=1, mode='auto')
callbacks = [es]
```

In [13]:

```
history_rmsprop = model.fit(x=x_train,
                           y=y_train,
                           batch_size=32,
                           epochs=5,
                           verbose=1,
                           callbacks=callbacks,
                           validation_split=0.25,
                           shuffle=True,
                           initial_epoch=0,
                           )
```

Train on 18750 samples, validate on 6250 samples

Epoch 1/5

```
18750/18750 [=====] - 171s 9ms/step - loss: 0.6735 - acc:
0.5806 - val_loss: 0.6449 - val_acc: 0.6312
```

Epoch 2/5

```
18750/18750 [=====] - 157s 8ms/step - loss: 0.6285 - acc:
0.6459 - val_loss: 0.5892 - val_acc: 0.6768
```

Epoch 3/5

```
18750/18750 [=====] - 157s 8ms/step - loss: 0.5684 - acc:
0.7079 - val_loss: 0.5238 - val_acc: 0.7426
```

Epoch 4/5

```
18750/18750 [=====] - 156s 8ms/step - loss: 0.5147 - acc:
0.7481 - val_loss: 0.5070 - val_acc: 0.7582
```

Epoch 5/5

```
18750/18750 [=====] - 157s 8ms/step - loss: 0.4734 - acc:
0.7770 - val_loss: 0.4430 - val_acc: 0.7939
```

In [14]:


```
import matplotlib.pyplot as plt

# 绘制训练 & 验证的准确率值
plt.plot(history_rmsprop.history['acc'])
plt.plot(history_rmsprop.history['val_acc'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()
```

<matplotlib.figure.Figure at 0x7f7563e6f208>

In [15]:

```
# 绘制训练 & 验证的损失值
plt.plot(history_rmsprop.history['loss'])
plt.plot(history_rmsprop.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()
```



In [16]:

```
base_model = VGG16(weights=None, include_top=False)
model_adam = build_model(base_model)
opt = keras.optimizers.Adam(lr=0.0001, decay=1e-6)
model_adam.compile(loss='categorical_crossentropy',
                   optimizer=opt,
                   metrics=['accuracy'])
```

<class 'keras.engine.training.Model'>

In [17]:

```

history_adam = model_adam.fit(x=x_train,
                              y=y_train,
                              batch_size=32,
                              epochs=5,
                              verbose=1,
                              callbacks=callbacks,
                              validation_split=0.25,
                              shuffle=True,
                              initial_epoch=0
                              )

```

Train on 18750 samples, validate on 6250 samples

Epoch 1/5

18750/18750 [=====] - 155s 8ms/step - loss: 0.6933 - acc: 0.4957 - val_loss: 0.6931 - val_acc: 0.5011

Epoch 2/5

18750/18750 [=====] - 154s 8ms/step - loss: 0.6930 - acc: 0.4991 - val_loss: 0.6931 - val_acc: 0.5021

Epoch 3/5

18750/18750 [=====] - 154s 8ms/step - loss: 0.6922 - acc: 0.5092 - val_loss: 0.6931 - val_acc: 0.4989

Epoch 4/5

18750/18750 [=====] - 154s 8ms/step - loss: 0.6749 - acc: 0.5645 - val_loss: 0.6523 - val_acc: 0.6051

Epoch 5/5

18750/18750 [=====] - 154s 8ms/step - loss: 0.6139 - acc: 0.6545 - val_loss: 0.5749 - val_acc: 0.6944


In [18]:

```

import matplotlib.pyplot as plt


# 绘制训练 & 验证的准确率值
plt.plot(history_adam.history['acc'])
plt.plot(history_adam.history['val_acc'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()

```



In [19]:

```
# 绘制训练 & 验证的损失值
plt.plot(history_adam.history['loss'])
plt.plot(history_adam.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()
```



In [20]:

```
base_model = VGG16(weights=None, include_top=False)
model_sgd = build_model(base_model)
opt = keras.optimizers.SGD(lr=0.0001, decay=1e-6)
model_sgd.compile(loss='categorical_crossentropy',
                  optimizer=opt,
                  metrics=['accuracy'])
```

<class 'keras.engine.training.Model'>

In [21]:

```

history_sgd = model_sgd.fit(x=x_train,
                             y=y_train,
                             batch_size=32,
                             epochs=5,
                             verbose=1,
                             callbacks=callbacks,
                             validation_split=0.25,
                             shuffle=True,
                             initial_epoch=0,
                             )

```

Train on 18750 samples, validate on 6250 samples

Epoch 1/5

18750/18750 [=====] - 154s 8ms/step - loss: 0.6931 - acc: 0.4996 - val_loss: 0.6931 - val_acc: 0.5011

Epoch 2/5

18750/18750 [=====] - 153s 8ms/step - loss: 0.6931 - acc: 0.5243 - val_loss: 0.6931 - val_acc: 0.5011

Epoch 3/5

18750/18750 [=====] - 154s 8ms/step - loss: 0.6931 - acc: 0.5099 - val_loss: 0.6931 - val_acc: 0.5013

Epoch 4/5

18750/18750 [=====] - 155s 8ms/step - loss: 0.6931 - acc: 0.5074 - val_loss: 0.6931 - val_acc: 0.5016

Epoch 5/5

18750/18750 [=====] - 155s 8ms/step - loss: 0.6931 - acc: 0.5023 - val_loss: 0.6931 - val_acc: 0.5021


In [22]:

```

import matplotlib.pyplot as plt


# 绘制训练 & 验证的准确率值
plt.plot(history_sgd.history['acc'])
plt.plot(history_sgd.history['val_acc'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()

```



In [23]:

```
# 绘制训练 & 验证的损失值
plt.plot(history_sgd.history['loss'])
plt.plot(history_sgd.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()
```



In [24]:

```
base_model = VGG16(weights=None, include_top=False)
model_large_lr = build_model(base_model)
opt = keras.optimizers.Adam(lr=0.001, decay=1e-6)
model_large_lr.compile(loss='categorical_crossentropy',
                        optimizer=opt,
                        metrics=['accuracy'])
```

<class 'keras.engine.training.Model'>

In [25]:

```

history_large_lr = model_large_lr.fit(x=x_train,
                                       y=y_train,
                                       batch_size=32,
                                       epochs=5,
                                       verbose=1,
                                       callbacks=callbacks,
                                       validation_split=0.25,
                                       shuffle=True,
                                       initial_epoch=0,
                                       )

```

Train on 18750 samples, validate on 6250 samples

Epoch 1/5

18750/18750 [=====] - 158s 8ms/step - loss: 0.6932 - acc: 0.5011 - val_loss: 0.6932 - val_acc: 0.4989

Epoch 2/5

18750/18750 [=====] - 157s 8ms/step - loss: 0.6932 - acc: 0.4956 - val_loss: 0.6932 - val_acc: 0.4989

Epoch 3/5

18750/18750 [=====] - 158s 8ms/step - loss: 0.6932 - acc: 0.4985 - val_loss: 0.6932 - val_acc: 0.4989

Epoch 4/5

18750/18750 [=====] - 157s 8ms/step - loss: 0.6932 - acc: 0.4930 - val_loss: 0.6932 - val_acc: 0.4989

Epoch 5/5


18750/18750 [=====] - 156s 8ms/step - loss: 0.6932 - acc: 0.4937 - val_loss: 0.6931 - val_acc: 0.5011

In [26]:

```


# 绘制训练 & 验证的准确率值
plt.plot(history_large_lr.history['acc'])
plt.plot(history_large_lr.history['val_acc'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()

```



In [27]:

```
plt.plot(history_large_lr.history['loss'])
plt.plot(history_large_lr.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()
```



In []: